

# Computational statistics, lecture 1

Frank Miller, Department of Computer and Information Science,  
Linköping University

[frank.miller@liu.se](mailto:frank.miller@liu.se)

October 31, 2023

# Course schedule

- Lecture 1: **Unidimensional optimization, computer arithmetic**
- Math-lecture 1: Basic matrix algebra, analytical optimization, determinants
- Lecture 2: **Multidimensional optimization**
- Math-lecture 2: Density, cumulative distribution function, integration
- Lecture 3: **Random number generation**
- Lecture 4: **Monte Carlo methods**
- Lecture 5: **Model selection and hypothesis testing**
- Lecture 6: **EM algorithm, stochastic optimization**

Teaching group: Krzysztof Bartoszek, L5-L6; Frank Miller, examiner, L1-L4, LM1-LM2; Bayu Brahmantio, Héctor Rodríguez Déniz, teaching assistants

Course homepage: <https://www.ida.liu.se/~732A90/index.en.shtml>; includes schedule, reading material, lecture notes, assignments

Computer labs: For each lecture; exercises to hand-in in **groups of 2**

# Evaluation of last course (HT2022)

- 23 students of around 50 submitted the evaluation; average grade 1.65
- Changes to HT2023:
  - Changes in content by deepening central topics (e.g. optimization) and instead removing some topics (e.g. reduced computer arithmetics); lab questions adjusted accordingly
  - Givens and Hoeting textbook slightly more used as basis for course layout (however, Gentle textbook covers topics as well)
  - New lecture slides for large parts of the course
  - Mathematical lectures moved inside the course (after L1 and after L2)

# Computational statistics

- When large or huge datasets should be analyzed and/or complex models are used, **statistics depends on effective computational methods**
- We will **learn in this course several algorithms** for optimization, randomization, Monte Carlo integration and **methods to use them**

# Today's content

- Optimization
  - Why?
  - Analytic univariate optimisation
  - Bi-section, Newton, and secant methods (univariate)
  - On convergence speed
- Computer arithmetics

# Optimization in statistics

- Maximum Likelihood
  - Minimizing risk in (Bayesian) decision theory
  - Minimizing sum of squares (Least Squares Estimate)
  - Maximizing information in experimental design
  - Machine learning
- Common problem in these examples:
    - $\mathbf{x}$   $p$ -dimensional vector,  $g: \mathbb{R}^p \rightarrow \mathbb{R}$  function
    - We search  $\mathbf{x}^*$  with  $g(\mathbf{x}^*) = \max g(\mathbf{x})$
- Minimization problem turns into maximization by considering  $-g$

# Least squares estimation (LSE)

- We search a Least Squares estimate  $\hat{\boldsymbol{\beta}}$  for  $\boldsymbol{\beta}$  minimising the distance  $g(\hat{\boldsymbol{\beta}}) = \|\hat{\mathbf{y}} - \mathbf{y}\|^2$  from  $\hat{\mathbf{y}} = \mathbf{X}\hat{\boldsymbol{\beta}}$  to  $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}$
- $g(\hat{\boldsymbol{\beta}}) = \|\mathbf{X}\hat{\boldsymbol{\beta}} - \mathbf{y}\|^2 = (\mathbf{X}\hat{\boldsymbol{\beta}} - \mathbf{y})^T (\mathbf{X}\hat{\boldsymbol{\beta}} - \mathbf{y}) = \hat{\boldsymbol{\beta}}^T \mathbf{X}^T \mathbf{X} \hat{\boldsymbol{\beta}} - 2\hat{\boldsymbol{\beta}}^T \mathbf{X}^T \mathbf{y} + \mathbf{y}^T \mathbf{y}$
- Setting the derivative to 0 ( $\frac{\partial f}{\partial \hat{\boldsymbol{\beta}}} = 2\mathbf{X}^T \mathbf{X} \hat{\boldsymbol{\beta}} - 2\mathbf{X}^T \mathbf{y} = 0$ ), we get  $\hat{\boldsymbol{\beta}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
- Optimization problem:
  - $\hat{\boldsymbol{\beta}}$   $p$ -dimensional vector,  $g: \mathbb{R}^p \rightarrow \mathbb{R}$  function
  - We search  $\hat{\boldsymbol{\beta}}$  with  $g(\hat{\boldsymbol{\beta}}) = \min g(\mathbf{b})$
- Here, we do not need to iteratively compute this minimum since we have an algebraic solution  $\hat{\boldsymbol{\beta}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

# Variations of least squares estimation

- Algebraic solution exists for the LSE, but not if we vary the problem
- Lasso estimate:  $g(\hat{\boldsymbol{\beta}}) = \|\mathbf{X}\hat{\boldsymbol{\beta}} - \mathbf{y}\|^2 + \lambda\|\hat{\boldsymbol{\beta}}\|_1$
- $L_1$ -estimation:  $g(\hat{\boldsymbol{\beta}}) = \|\mathbf{X}\hat{\boldsymbol{\beta}} - \mathbf{y}\|_1$
- Many further variations of estimates have been considered
- In all cases, we search  $\hat{\boldsymbol{\beta}}$  with  $g(\hat{\boldsymbol{\beta}}) = \min g(\mathbf{b})$
- Recall: Norms for  $\mathbf{x} = (x_1, \dots, x_p)^T$ :  $\|\mathbf{x}\| = \|\mathbf{x}\|_2 = \sqrt{x_1^2 + \dots + x_p^2}$  (Euclid),  $\|\mathbf{x}\|_1 = |x_1| + \dots + |x_p|$ ,  $\|\mathbf{x}\|_\infty = \max\{|x_1|, \dots, |x_p|\}$  (max-norm)



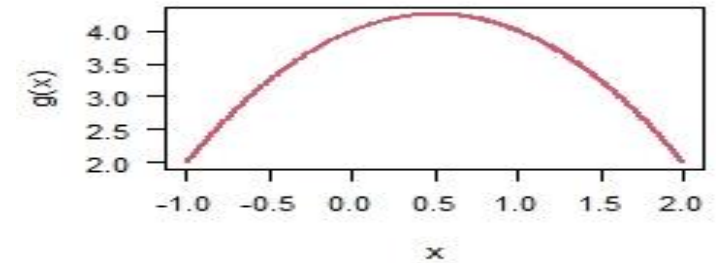


# Univariate optimization

- $x$  real number,  $g: \mathbb{R} \rightarrow \mathbb{R}$  continuously differentiable function
- We search  $x^*$  with  $g(x^*) = \max g(x)$
- Compute  $g'(x)$  and search  $x^*$  with  $g'(x^*) = 0$
- One has then to check if the result is maximum, minimum, possibly local optimum...


# Univariate optimization – analytic solution

- Compute  $g'(x)$  and search  $x^*$  with  $g'(x^*) = 0$
- Example, where analytic optimization possible:
  - $g(x) = 4 + x - x^2$   
 $g'(x) = 1 - 2x$   
 $g'(x) = 0$  if and only if  $x = 1/2$   
 $g''(x) = -2, \quad g''(1/2) = -2 < 0$
- Therefore,  $g$  has local maximum at  $x = 1/2$
- Now, we have cases in mind where the **analytic solution is not possible**, and we need **iterative methods**



# Univariate optimization: bisection

- Search  $x^*$  with  $g'(x^*) = 0$ :
  - 1) Start with interval  $[a_0, b_0]$  such that  $g'(a_0) \cdot g'(b_0) < 0, t = 0$
  - 2) Set  $x^{(t)} = (a_t + b_t)/2$
  - 3) Define next interval  $[a_t, b_t]$  by  
 $[a_t, x^{(t)}]$  if  $g'(a_t) \cdot g'(x^{(t)}) \leq 0$ ,  
 $[x^{(t)}, b_t]$  if  $g'(x^{(t)}) \cdot g'(b_t) < 0$
  - 4) Set  $t$  to  $t+1$  and go to 2)
- See [video on course homepage](#)
- **Iteratively** improve approximations for  $x^*$ :  
 $x^{(0)} \rightarrow x^{(1)} \rightarrow x^{(2)} \rightarrow \dots$



# Optimization: convergence criterion

- Compare  $x^{(t)}$  and  $x^{(t+1)}$  and stop if they are “close enough”
- Absolute convergence criterion:

$$|x^{(t+1)} - x^{(t)}| < \epsilon$$

- Relative convergence criterion:

$$\frac{|x^{(t+1)} - x^{(t)}|}{|x^{(t)}|} < \epsilon$$

# Univariate Newton(-Raphson)

- $x$  real number,  $g: \mathbb{R} \rightarrow \mathbb{R}$  twice differentiable function
- Search  $x^*$  with  $g(x^*) = \max g(x)$  by searching  $x^*$  with  $g'(x^*) = 0$

- Taylor expansion around  $x^*$  motivates:

$$0 = g'(x^*) \approx g'(x^{(t)}) + (x^* - x^{(t)})g''(x^{(t)})$$

$$-(x^* - x^{(t)})g''(x^{(t)}) \approx g'(x^{(t)})$$

$$x^* \approx x^{(t)} - g'(x^{(t)})/g''(x^{(t)})$$

- Therefore, the Newton-iteration works as:

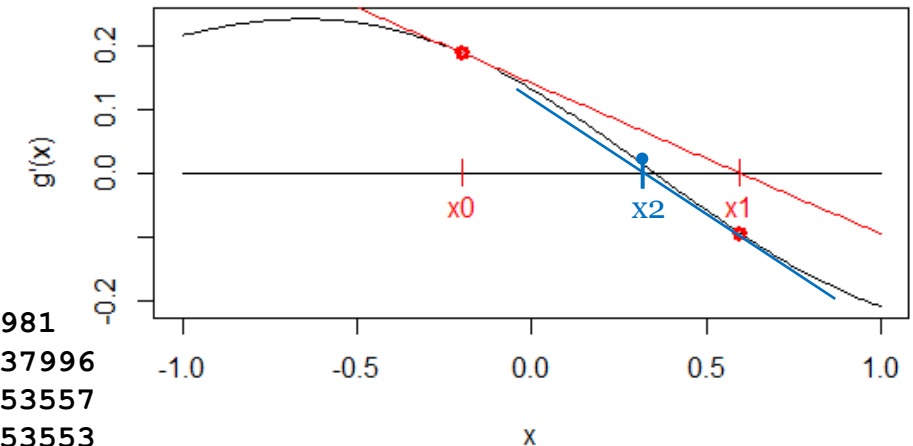
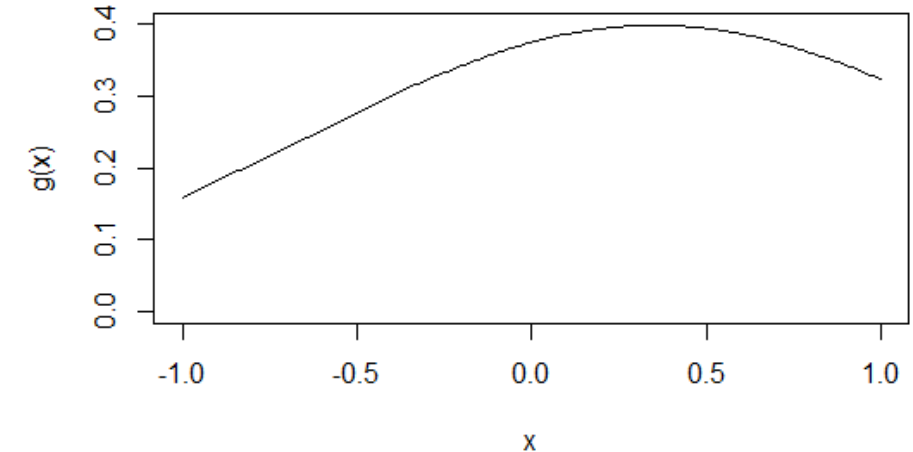
$$x^{(t+1)} = x^{(t)} - g'(x^{(t)})/g''(x^{(t)})$$

# Univariate Newton(-Raphson)

- $x^{(t+1)} = x^{(t)} - g'(x^{(t)})/g''(x^{(t)})$
- Start with a  $x^{(0)}$
- Tangent in  $(x^{(0)}, g'(x^{(0)}))$  determines  $x^{(1)}$
- Tangent in  $(x^{(1)}, g'(x^{(1)}))$  determines  $x^{(2)}$
- ...
- until convergence criterion met

+Newton method is fast

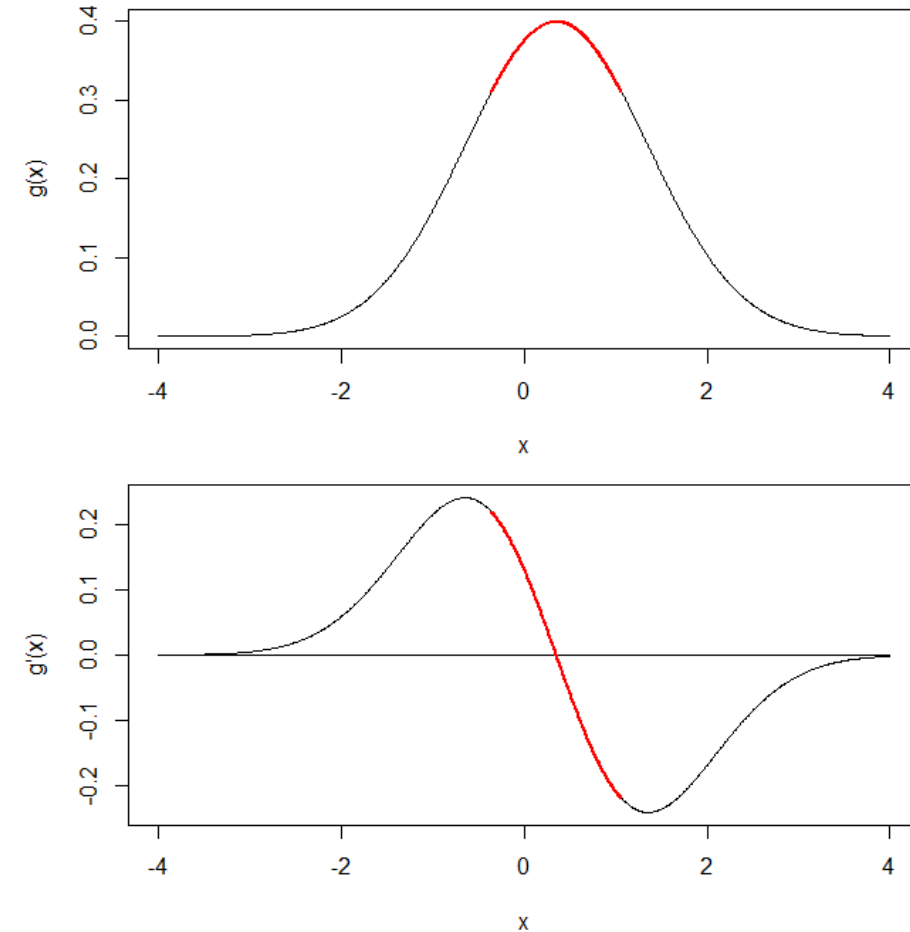
- Requires existence and computation of  $g''$



```
x0 -0.2
x1 0.5981
x2 0.337996
x3 0.353557
x4 0.353553
x5 0.353553
STOP
```

# Univariate Newton(-Raphson)

- $x^{(t+1)} = x^{(t)} - g'(x^{(t)})/g''(x^{(t)})$
- What about the starting value  $x^{(0)}$ ?



# Univariate secant method

- $x$  real number,  $g: \mathbb{R} \rightarrow \mathbb{R}$  once differentiable function
- Search  $x^*$  with  $g(x^*) = \max g(x)$  by searching  $x^*$  with  $g'(x^*) = 0$
- Recall: The Newton-iteration works as:
$$x^{(t+1)} = x^{(t)} - g'(x^{(t)})/g''(x^{(t)})$$
- Need to compute  $g''$  which might be difficult. Instead:
- Approximate  $g''(x^{(t)})$  by  $[g'(x^{(t)}) - g'(x^{(t-1)})]/(x^{(t)} - x^{(t-1)})$



# Univariate secant method

- $x^{(t+1)} = x^{(t)} - g'(x^{(t)}) \frac{x^{(t)} - x^{(t-1)}}{g'(x^{(t)}) - g'(x^{(t-1)})}$
- Start with  $x^{(0)}$  and  $x^{(-1)}$
- Secant through  $x^{(0)}$  and  $x^{(-1)}$  determines  $x^{(1)}$
- Secant through  $x^{(1)}$  and  $x^{(0)}$  determines  $x^{(2)}$
- ...
- until stopping crit. fulfilled

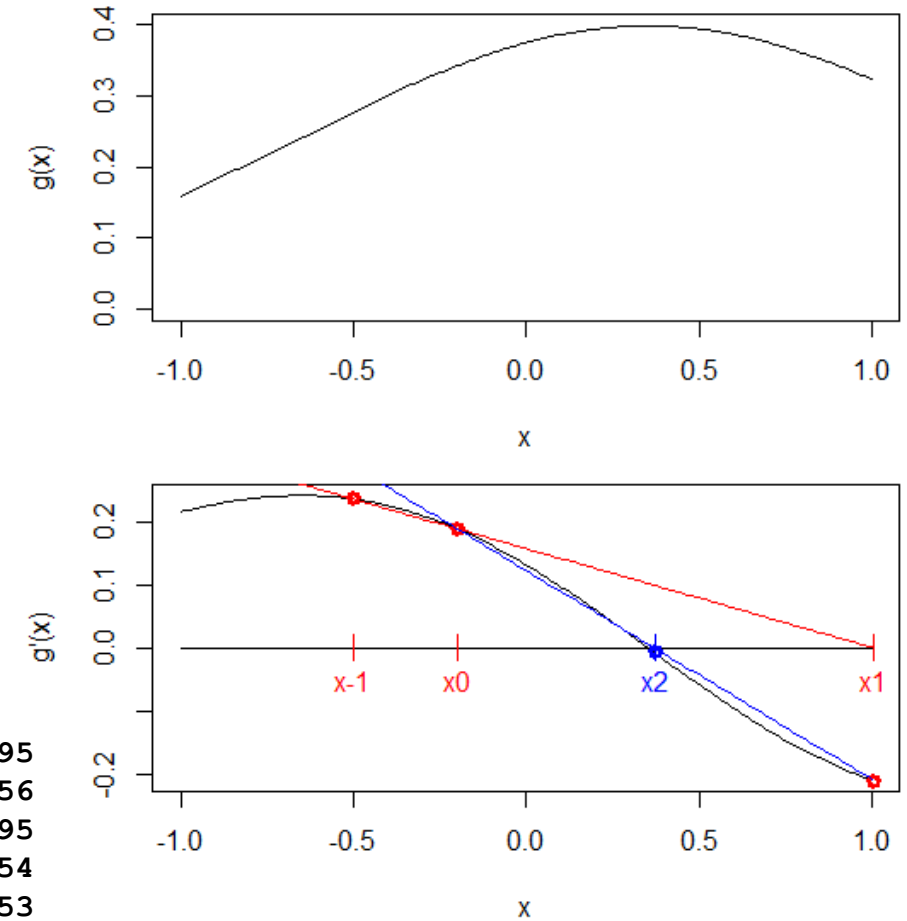
• Quite fast

• No 2<sup>nd</sup> derivative necessary

```

x0 -0.2
x1 1.006995
x2 0.371656
x3 0.349095
x4 0.353554
x5 0.353553
x6 0.353553
STOP

```



# Convergence speed of optimization algorithms

- Convergence speed can be quantified by  $q$  and  $c$  as follows:
  - Let  $\varepsilon^{(t)} = x^{(t)} - x^*$ ,
  - Find  $q$  and  $c$  such that  $\lim_{t \rightarrow \infty} \varepsilon^{(t+1)} / (\varepsilon^{(t)})^q = c$
- $\varepsilon = 1, 0.5, 0.25, 0.125, 0.063, 0.031, \dots \rightarrow q=1, c=0.5,$
- $\varepsilon = 1, 0.1, 0.01, 0.001, 0.0001, \dots \rightarrow q=1, c=0.1,$
- If  $q=1$ , we say that convergence is "linear"
- $\varepsilon = 1, 0.5, 0.125, 0.008, 0.00003, \dots \rightarrow q=2, c=0.5.$
- If  $q=2$ , we say that convergence is "quadratic"

Convergence  
order

Convergence  
rate

Intuitively,  
 $\varepsilon^{(t+1)} \approx c \cdot (\varepsilon^{(t)})^q$

# Determine empirically convergence rate (and order) of optimization algorithms

- You have a given optimization algorithm and you have determined or know the maximiser  $\mathbf{x}^*$ . To check convergence speed in an optimization-run, you can calculate

$$D^{(t)} = \frac{|x^{(t)} - x^*|}{|x^{(t-1)} - x^*|}$$

(see Givens and Hoeting, 2013, page 101/102, for an example)

- If  $D^{(t)} \rightarrow 1$ , there is not even linear convergence (bad, order  $q < 1$ ),  
If  $D^{(t)} \rightarrow c \in (0,1)$ , linear convergence (order  $q=1$ ) with rate  $c$ ,  
If  $D^{(t)} \rightarrow 0$ , better than linear convergence (order  $q > 1$ ).

# Comparison of univariate optimization methods

Bisection	Secant	Newton
$g'$ required	$g'$ required	$g''$ required
finds always an optimum between $a_0$ and $b_0$ (but could be local)	converges only when the two starting values "close" to optimum	converges only when starting value "close" to optimum
slow $q=1$	$q = \frac{1 + \sqrt{5}}{2} = 1.62$	fast $q=2$

- There are also algorithms not needing  $g'$
- R-function **optimize** uses such an algorithm ( $q=1.324$ )

# Computer arithmetics

- Numbers are represented as binary numbers ( $17 = 1 * 2^4 + 1 * 2^1 = "1001"$ )
- Rational numbers are also represented based on the binary system:  

$$\pm 0.d_1 d_2 \dots d_p * 2^e,$$

$$e = \pm e_1 e_2 \dots e_q$$
- E.g.  $p = 52, q = 10$ , two signs  $\Rightarrow$  one number needs 64 bits in the computer
- Limits in representation depending on  $p$  and  $q$

```
> 2^1023
[1] 8.988466e+307
> 2^1024
[1] Inf
(overflow)
```

```
> 2^-1074
[1] 4.940656e-324
> 2^-1075
[1] 0
(underflow)
```

```
> 3/5-2/5-1/5
[1] -5.551115e-17
> if (3/5-2/5==1/5)
>   print("yes") else
>   print("no")
[1] "no"
```

# Computer arithmetics

- Good to have limitations of computer arithmetics in mind!
- Example: Binomial coefficient (avoid overflow)

$$\binom{n}{k} = \frac{n!}{k!(n-k)!} = \frac{(k+1)(k+2)\cdots(n-1)n}{(n-k)!}$$

$$\binom{200}{2} = \frac{200!}{2!198!} = \frac{3*4*\cdots*199*200}{1*2*\cdots*197*198} = \frac{199*200}{1*2} = 19900$$

```
> n <- 200
> k <- 2
> prod(1:n) / (prod(1:k)*prod(1:(n-k)))
[1] NaN
> prod((k+1):n) / (1:(n-k))
[1] 19900
```

# Course material, lab, seminar, exam

- Homepage: <https://www.ida.liu.se/~732A90/index.en.shtml>
  - Lecture notes, lab- and seminar info, exam info
- Submission of 6 labs via LISAM – all need to be passed – groups of 2
  - First lab: Oct 31 to Nov 7
- Mandatory attendance at 3 seminars and 1 presentation or opposition
- Computer exam: Jan 9, 2024. Own document of 100 pages can be used.
- 10 points to pass (E); 12 or more: D;  $\geq 14$ : C;  $\geq 16$ : B;  $\geq 18$ : A

## Literature:

- Gentle JE (2009). *Computational Statistics*, Springer
- Givens GH, Hoeting JA (2013). *Computational Statistics*, Wiley